

Literature Review & Elaboration of Problem

Machine Learning vs. Deep Learning for

Social Media-Based Mental Health

Detection

Version: 2.1

Project Code	F25-112
Internal Supervisor	Ms. Fizza Aqeel
External Supervisor	—
Project Manager	—
Project Team	Muhammad Tahir Haider Muhammad Huzaifa Saleem Syed Muhammad Sufyan
Submission Date	December 9, 2025

Document History

Version	Name of Person	Date	Description of change
1.0	Tahir	14-Nov-25	Document created, Added References and Bibliography
1.1	Sufyan	15-Nov-25	Added Abstract
1.2	Huzaiifa	20-Nov-25	Added Background and Justification
1.3	Huzaiifa	21-Nov-25	Updated Background and Justification, Added Problem Statement
1.4	Sufyan	23-Nov-25	Added Literature Review
1.5	Tahir	27-Nov-25	Updated Literature Review
1.6	Tahir	28-Nov-25	Updated Literature Review
1.7	Huzaiifa	29-Nov-25	Updated Literature Review
1.8	Sufyan	30-Nov-25	Updated Literature Review, Updated References and Bibliography
1.9	Tahir	1-Dec-25	Added Appendices
2.0	Sufyan	3-Dec-25	Updated Appendices and Table of Contents
2.1	Tahir	5-Dec-25	Final review and document preparation for submission

[illegible]

Table of Contents

1. Abstract.....	5
2. Background and Justification.....	5
3. Problem Statement	5
3.1. <i>Research Questions</i>	6
3.2. <i>Challenges Anticipated</i>	6
4. Literature Review	6
4.1. <i>The Critical Need for Early Detection</i>	6
4.2. <i>Social Media Data: Sources, Characteristics, and Challenges</i>	6
4.2.1. <i>Dominant Data Sources</i>	6
4.2.2. <i>Preprocessing and Feature Engineering</i>	7
4.3. <i>Modeling Approaches for Mental Health Detection</i>	7
4.3.1. <i>Traditional Machine Learning Models</i>	7
4.3.2. <i>Deep Learning Models</i>	7
4.3.3. <i>Transformer-Based Models</i>	8
4.3.4. <i>Existing Comparative Studies</i>	8
4.4. <i>Challenges, Limitations, and Research Gaps</i>	8
4.4.1. <i>Lack of Standardized Benchmarking</i>	8
4.4.2. <i>Multi-Condition Detection Challenges</i>	8
4.4.3. <i>Computational Feasibility and Scalability</i>	9
4.4.4. <i>Dataset Limitations and Evaluation Inconsistencies</i>	9
5. Appendices.....	10
<i>Appendix A: Summary of Model Characteristics</i>	10
<i>Appendix B: Performance Evaluation Metrics</i>	11
<i>Appendix C: Glossary of Key Terms</i>	12
6. References	14
7. Bibliography.....	16

1. Abstract

The escalating global prevalence of mental health disorders, such as depression and anxiety, underscores an urgent need for innovative and scalable detection mechanisms to enable early intervention and appropriate treatment. Traditional diagnostic methods often face challenges related to stigma, lack of awareness, and complexity in identifying the disorders before they escalate into severe crises, including suicidal ideation. Given the widespread integration of social media platforms into daily life, these sites have become rich, real-time sources of user-generated content reflecting emotional states and psychological distress. This project utilizes NLP techniques to analyze social media text for the automated and early detection of depression, and suicidal ideation. The methodology encompasses several crucial steps: systematic data preprocessing to handle the informal nature of social media language; feature extraction using various Natural Language Processing (NLP) techniques; and comprehensive model training and evaluation. We employ and comparatively evaluate a spectrum of Machine Learning (ML) models—Logistic Regression (LR), Support Vector Machine (SVM), Naïve Bayes (NB), and Random Forest (RF)—alongside advanced Deep Learning (DL) architectures, including Long Short-Term Memory (LSTM), Bidirectional LSTM (BiLSTM), Convolutional Neural Networks (CNN), and Bidirectional Encoder Representations from Transformers (BERT). The significance of this research lies in its potential to complement conventional clinical assessments by providing scalable, high-accuracy tools for predicting mental health risks including depression, and suicidal ideation, thereby contributing significantly to timely intervention strategies and improved global mental health outcomes.

2. Background and Justification

The escalating prevalence of mental health disorders worldwide necessitates scalable, automated detection methods to supplement traditional clinical assessments. While clinical diagnosis remains the gold standard, the widespread adoption of social media platforms has created unprecedented opportunities for early detection through analysis of user-generated content [1], [2].

This project builds upon a substantial body of research demonstrating the feasibility of detecting mental health conditions from social media text. However, the literature reveals inconsistencies in methodological approaches, with studies employing varying preprocessing techniques, feature extraction methods, and model architectures, making cross-study comparisons difficult [1].

This research is justified by the need for a systematic, controlled comparison of classical machine learning, deep learning, and transformer-based approaches using consistent datasets and evaluation frameworks. By utilizing established benchmark datasets for depression, and suicidal ideation detection, this project provides reproducible results that can guide future research and practical deployment decisions.

The inclusion of suicidal ideation detection alongside depression addresses a critical gap, as early identification of suicidal thoughts can enable life-saving interventions. The comparative nature of this research will reveal which modeling approaches offer the best balance of accuracy, interpretability, and computational efficiency for real-world mental health surveillance applications.

3. Problem Statement

Despite significant research progress in mental health detection from social media, several critical problems remain:

- 1. Lack of Standardized Benchmarking:** *Existing studies employ heterogeneous datasets, preprocessing pipelines, and evaluation protocols, making it difficult to determine which modeling approaches are genuinely superior versus simply benefiting from better data or preprocessing [1], [2].*
- 2. Interpretability-Accuracy Trade-off:** *While transformer models achieve state-of-the-art accuracy, their "black box" nature limits clinical interpretability. Conversely, classical ML models offer transparency but may sacrifice performance [2]. The optimal balance for practical deployment remains unclear.*
- 3. Computational Feasibility:** *Large-scale deployment of mental health screening systems requires consideration of computational costs. The relative efficiency of different model families in achieving acceptable performance thresholds has not been systematically quantified.*

- 4. Multi-condition Detection:** Real-world mental health scenarios involve comorbid conditions. Most studies focus on single-condition classification, but practical systems must detect multiple co-occurring conditions such as depression, and suicidal ideation simultaneously.

3.1. Research Questions

- Which model family provides the best performance for detecting depression, and suicidal ideation from social media text?
- What are the quantifiable trade-offs between model accuracy, interpretability, and computational requirements?
- Do different conditions benefit from different modeling approaches?

3.2. Challenges Anticipated

- Handling class imbalance in mental health datasets.
- Ensuring fair comparison across diverse model architectures.
- Managing computational resources for training large transformer models.
- Interpreting and validating model predictions without clinical expertise.

4. Literature Review

This section provides a detailed review and synthesis of the scholarly literature pertaining to the detection of mental health disorders from social media text, focusing on the intersection of NLP and ML techniques. The purpose is to provide a clear understanding of the research problem, demonstrate the authenticity and importance of the proposed work, identify current methodological limitations, and contextualize the technical approach adopted by this project.

4.1. The Critical Need for Early Detection

The detection of mental health issues, particularly depression and anxiety, is a global health priority, with depression affecting over 280 million people worldwide [1]. Clinicians often face difficulty in identifying these mental disorders due to their complicated and subjective nature, leading to potential misdiagnosis or delayed treatment. This delay can exacerbate symptoms, sometimes resulting in catastrophic outcomes such as suicide [3].

Social media platforms have emerged as invaluable, real-time data reservoirs, capturing user-generated content that reflects emotional and psychological states [2]. The observation that subjects with mental disorders exhibit particular characteristics in their social media usage patterns confirms the feasibility of early, automated text-based detection [4]. Beyond depression and anxiety, the detection of suicidal ideation represents a particularly critical application, as early identification of suicidal thoughts can enable life-saving interventions [5].

The integration of computational approaches with traditional clinical assessment offers potential for scalable, non-invasive detection methods that can complement existing healthcare infrastructure [2]. However, the development of such systems requires careful consideration of accuracy, interpretability, and ethical implications.

4.2. Social Media Data: Sources, Characteristics, and Challenges

4.2.1. Dominant Data Sources

The choice of social media platform significantly influences the linguistic characteristics of the dataset analyzed. Twitter (X) is the most widely utilized platform for mental health research, accounting for over 60% of studies in this area, primarily due to the platform's accessibility for data collection [1], [5]. Reddit is also extensively used, especially in subreddits related to mental health (e.g., r/depression), offering longer, more descriptive posts conducive to detailed linguistic analysis [6], [7]. Data collection methods frequently rely on searching public content using keywords, phrases, or regular expressions related to self-reported diagnoses (e.g., "I was diagnosed with depression") [8].

4.2.2. Preprocessing and Feature Engineering

Converting raw, unstructured social media text into features suitable for machine learning models represents a challenging but vital step [5]. Standard preprocessing techniques include converting text to lowercase, removing URLs, special characters, and stop words [5], [9]. However, the informal nature of social media text introduces significant noise, including slang, emojis, and implicit expressions like sarcasm. Advanced preprocessing techniques are required to normalize slang through customized dictionaries and interpret emojis by converting them into text or calculating associated sentiment scores [10].

A particularly significant methodological challenge is the inadequate handling of negations, which can severely skew sentiment analysis if overlooked. Studies indicate that only a minority of research (around 23%) explicitly addresses negation handling [1], representing a critical gap in preprocessing pipelines.

For feature representation, researchers employ both lexical features such as N-grams and Term Frequency-Inverse Document Frequency (TF-IDF), as well as psycholinguistic features extracted using dictionaries like Linguistic Inquiry and Word Count (LIWC) [2], [8], [11]. LIWC analysis frequently highlights features associated with depression, including high frequencies of first-person singular pronouns and negative emotion words [8], [12]. Additionally, quantified emotion scores derived from sentiment analysis tools like TextBlob or the NRC Emotion Lexicon (EmoLex) are often used to identify depression-related emotional cues [8], [10], [11].

4.3. Modeling Approaches for Mental Health Detection

Research in this domain spans traditional machine learning models, deep learning architectures, and transformer-based approaches. This section reviews each category and their comparative performance characteristics.

4.3.1. Traditional Machine Learning Models

Classical machine learning models remain valuable due to their computational efficiency and high degree of interpretability compared to deep learning approaches, particularly for moderately sized datasets [2].

Support Vector Machine (SVM) is among the most widely reported ML classifiers for text classification tasks, often employing different kernels to effectively separate classes [1], [8], [13], [14]. Hybrid models integrating SVM, such as the NeuroVibeNet framework, have achieved high accuracy in multimodal detection scenarios [3].

Logistic Regression (LR) is particularly favored for its interpretability, as it readily captures positive and negative feature associations relevant to clinical interpretation [2]. When effectively tuned, LR has demonstrated accuracy comparable to SVM in distinguishing depressed users from control groups [5], [15].

Naïve Bayes (NB) is primarily employed as a simple, efficient baseline model due to its probabilistic nature and computational efficiency [8], [14].

Random Forest (RF) is an ensemble method that consistently yields robust performance in depression detection tasks through its use of multiple decision trees [8], [15], [16]. Advanced tree-based models like XGBoost have demonstrated exceptional performance in multi-level severity classification [17].

4.3.2. Deep Learning Models

Deep learning models excel by automatically learning hierarchical, complex feature representations, bypassing the constraints of manual feature engineering required by traditional ML methods [2], [9].

Convolutional Neural Networks (CNN) are employed to capture localized, high-level features and hierarchical representations within social media posts. CNNs have been successfully combined with linguistic metadata for effective early detection, and hybrid architectures like SBERT-CNN (combining Sentence BERT embeddings with CNN layers) have achieved impressive F1 scores of 0.86 in depression identification [9], [13].

Long Short-Term Memory (LSTM) and Bidirectional LSTM (BiLSTM) networks are fundamental for processing sequential text data. LSTMs effectively capture long-range dependencies, while BiLSTM surpasses standard LSTM performance by processing input context in both forward and backward directions, thereby improving sequential pattern detection [10]. BiLSTM models augmented with attention mechanisms have demonstrated particularly competitive results in text-based depression detection [4].

4.3.3. Transformer-Based Models

Transformer architectures, particularly BERT (Bidirectional Encoder Representations from Transformers), leverage bidirectional encoding and attention mechanisms to create highly contextualized word representations, establishing state-of-the-art benchmarks for accuracy in complex NLP tasks [5], [7], [10]. BERT's bidirectional attention mechanism enables it to capture contextual nuances critical for mental health detection, such as understanding how negations or qualifiers modify emotional expressions. Unlike traditional models that process text sequentially, BERT considers both left and right context simultaneously, making it particularly effective at understanding subtle linguistic patterns associated with mental health conditions [5], [14], [18].

Fine-tuned variants of BERT demonstrate superior classification performance in depression detection. Domain-specific models like MentalBERT, pre-trained on mental health corpora, have achieved the highest reported F1-scores (97.31%) in multi-level depression severity detection, underscoring the value of specialized pre-training [17]. Alternative variants like RoBERTa and DistilBERT offer performance-efficiency trade-offs suitable for resource-constrained scenarios [5], [14], [18].

To maximize performance, BERT embeddings are often combined with recurrent architectures. The BERT-BiLSTM hybrid, which integrates BERT's contextual understanding with BiLSTM's sequential modeling capabilities, has demonstrated enhanced performance in capturing complete depression-related content, often yielding superior recall and F1 scores compared to using either architecture alone [4], [10].

4.3.4. Existing Comparative Studies

Several recent studies have attempted systematic comparisons between modeling approaches, but with varying methodologies and conclusions. Research comparing ML and DL approaches has found that while transformer models achieve highest accuracy, classical ML remains competitive when interpretability or computational resources are prioritized [2], [5], [14]. Studies examining performance trade-offs note that deep learning excels in accuracy while traditional ML offers advantages in efficiency and transparency [2]. However, these comparative studies reveal methodological limitations including inconsistent preprocessing pipelines across model families, evaluation on different datasets or data splits, and predominant focus on binary classification rather than multi-class detection [1], [5]. Furthermore, few studies provide detailed analysis of computational costs alongside performance metrics, limiting practical guidance for deployment decisions [2].

4.4. Challenges, Limitations, and Research Gaps

Despite significant advances, the literature reveals critical gaps that warrant further investigation:

4.4.1. Lack of Standardized Benchmarking

Existing studies employ heterogeneous datasets, preprocessing pipelines, and evaluation protocols, making it difficult to determine which modeling approaches are genuinely superior versus simply benefiting from better data or preprocessing decisions [1], [2]. This lack of standardization means that reported performance differences may reflect variations in data quality, preprocessing choices, or evaluation metrics rather than true algorithmic superiority.

Moreover, the interpretability-accuracy trade-off remains poorly quantified across studies. While transformer models achieve state-of-the-art accuracy, their "black box" nature limits clinical interpretability and trust. Conversely, classical ML models offer transparency but may sacrifice performance [2]. Without standardized comparison frameworks using consistent datasets and evaluation protocols, determining the optimal balance for practical deployment remains unclear.

4.4.2. Multi-Condition Detection Challenges

A significant limitation in existing research is the predominant focus on binary classification (e.g., depressed vs. non-depressed), which fails to reflect the complexity of real-world mental health scenarios where individuals often experience comorbid conditions [16]. Recent work has begun addressing multi-class classification, attempting to simultaneously distinguish between depression, anxiety, and normal mental states, but such studies remain limited [7], [16].

Multi-condition detection systems face unique challenges including overlap in linguistic markers between different disorders, the need for larger and more balanced datasets representing multiple conditions, and increased complexity in model architecture and training. Deep learning and transformer approaches have shown promise in capturing nuanced differences between conditions, but interpretability challenges become more pronounced in multi-class scenarios [17]. Critically, existing comparative studies rarely evaluate model performance across multiple mental health conditions simultaneously, making it difficult to determine which modeling approaches generalize best across different disorders.

4.4.3. Computational Feasibility and Scalability

While transformer models achieve superior accuracy, their computational requirements present significant barriers to real-world deployment. Training large models like BERT requires substantial GPU resources and time, which may not be feasible for all research contexts or real-time monitoring applications [2]. The literature reveals limited systematic analysis of performance-to-cost ratios across model families.

Classical ML models, despite potentially lower accuracy in some scenarios, offer clear advantages in training speed, inference time, and resource requirements. However, existing studies rarely provide quantitative comparisons of computational costs—including training time, memory usage, and inference latency—alongside performance metrics [2]. This gap is particularly relevant for scaling mental health detection systems to monitor large user populations in real-time, where computational efficiency becomes a critical practical constraint.

4.4.4. Dataset Limitations and Evaluation Inconsistencies

While community-labeled benchmark datasets available on platforms like Kaggle have democratized mental health detection research [1], [2], these datasets exhibit important limitations. Most are predominantly English-language content from specific platforms (primarily Twitter and Reddit), potentially limiting generalizability across languages, cultures, and platform contexts [1]. Sampling biases related to self-selection of users in mental health subreddits may not represent the broader population of individuals experiencing mental health issues.

The scarcity of clinically validated, annotated datasets remains a fundamental challenge, as most research relies on self-reported diagnoses or community labels rather than professional clinical assessment [1]. Additionally, evaluation practices show inconsistencies across studies, with different train-test splits, cross-validation strategies, and performance metrics further complicating comparative analysis [2].

5. Appendices

Appendix A: Summary of Model Characteristics

Table A.1: Comparison of Model Families for Mental Health Detection

Model Type	Examples	Key Strengths	Key Limitations	Typical Use Cases
Classical Machine Learning	Logistic Regression (LR) Support Vector Machine (SVM) Naïve Bayes (NB) Random Forest (RF)	High interpretability	Requires manual feature engineering	Baseline model establishment
		Computationally efficient	May have lower accuracy than deep learning	Resource-constrained scenarios
		Works well with smaller datasets	Limited ability to capture complex patterns	Clinical settings requiring interpretability
		Fast training and inference	Performance depends on feature quality	Real-time applications
		Clinically transparent		
Deep Learning	Convolutional Neural Networks (CNN) Long Short-Term Memory (LSTM) Bidirectional LSTM (BiLSTM)	Automatic feature learning	Requires substantial training data	When accuracy is priority
		Captures sequential patterns	Less interpretable ("black box")	Sufficient training data available
		Handles complex, hierarchical representations	Higher computational cost	Sequential pattern detection needed
		Better performance with larger datasets	Longer training time	Complex linguistic analysis
			Risk of overfitting	
Transformer-Based Models	BERT RoBERTa DistilBERT MentalBERT	State-of-the-art accuracy	Highest computational requirements	When highest accuracy is essential
		Captures contextual nuances bidirectionally	"Black box" nature	Computational resources available
		Pre-trained models available	Requires significant GPU resources	Fine-tuning with domain-specific data
		Transfer learning capabilities	Longest training time	
		Understands complex linguistic patterns	Large memory footprint	Research benchmarking

Note: This table synthesizes findings from the literature review and guides model selection and evaluation in this comparative study.

Appendix B: Performance Evaluation Metrics

This project evaluates models using the following standard classification metrics:

Primary Metrics

- 1. Accuracy:** The proportion of correct predictions (both true positives and true negatives) among the total number of cases examined.
Formula: $(TP + TN) / (TP + TN + FP + FN)$
- 2. Precision:** The proportion of true positive predictions among all positive predictions made by the model.
Formula: $TP / (TP + FP)$
- 3. Recall (Sensitivity):** The proportion of actual positive cases that were correctly identified by the model.
Formula: $TP / (TP + FN)$
- 4. F1-Score:** The harmonic mean of precision and recall, providing a single balanced measure that considers both false positives and false negatives.
Formula: $2 \times (Precision \times Recall) / (Precision + Recall)$
- 5. ROC-AUC (Receiver Operating Characteristic - Area Under Curve):** Measures the model's ability to distinguish between classes across all possible classification thresholds. AUC represents the area under the ROC curve.
Range: 0 to 1, where 1.0 indicates perfect classification and 0.5 indicates random guessing.

Notation Key

- **TP (True Positives):** Correctly identified positive cases (e.g., correctly detected mental health condition)
- **TN (True Negatives):** Correctly identified negative cases (e.g., correctly identified no mental health condition)
- **FP (False Positives):** Incorrectly identified positive cases (Type I error)
- **FN (False Negatives):** Incorrectly identified negative cases (Type II error)

Additional Considerations

Class Imbalance Handling: Mental health datasets often exhibit class imbalance, where the number of positive cases (individuals with mental health conditions) may be significantly smaller than negative cases. In such scenarios, accuracy alone can be misleading. F1-Score and ROC-AUC provide more reliable performance indicators for imbalanced datasets.

Multi-Class Classification: For distinguishing between depression, anxiety, and normal states simultaneously, metrics will be calculated using:

- **Macro-averaging:** Calculating metrics independently for each class and taking the average (treats all classes equally)
- **Weighted-averaging:** Calculating metrics for each class and taking the average weighted by class support (accounts for class imbalance)

Note: All models in this comparative study will be evaluated using these standardized metrics to ensure fair and consistent performance comparison.

Appendix C: Glossary of Key Terms

Anxiety: A mental health condition characterized by excessive worry, nervousness, and fear that interferes with daily activities. Includes various subtypes such as generalized anxiety disorder, social anxiety, and panic disorder.

Attention Mechanism: A technique that allows models to focus on different parts of the input sequence when making predictions, assigning importance weights to different tokens based on their relevance to the current context.

Baseline Model: A simple model used as a reference point for comparison, establishing minimum expected performance that more sophisticated models should exceed.

Benchmark Dataset: A standardized dataset used for evaluating and comparing the performance of different models, ensuring fair comparison by providing consistent evaluation conditions.

BERT (Bidirectional Encoder Representations from Transformers): A transformer-based language model developed by Google that processes text bidirectionally to create contextualized word representations. Pre-trained on large text corpora and fine-tunable for specific tasks.

BiLSTM (Bidirectional Long Short-Term Memory): A recurrent neural network architecture that processes sequences in both forward and backward directions simultaneously, capturing context from both past and future tokens.

CNN (Convolutional Neural Network): A deep learning architecture originally designed for image processing, adapted for text classification by applying convolutional filters to capture local patterns and hierarchical features in text.

Comorbidity: The simultaneous presence of two or more medical or mental health conditions in a patient, such as depression and anxiety occurring together. Common in mental health where multiple disorders frequently co-occur.

Cross-Validation: A resampling technique for assessing model performance and generalizability by partitioning data into complementary subsets,

training on some subsets and validating on others, then rotating the validation subset.

Depression: A common mental health disorder characterized by persistent sadness, loss of interest, changes in sleep and appetite, fatigue, and difficulty concentrating. Major depressive disorder affects over 280 million people globally.

DistilBERT: A distilled (compressed) version of BERT that retains 97% of BERT's language understanding while being 60% faster and having 40% fewer parameters, making it more efficient for deployment.

Embedding: A learned dense vector representation of words, sub words, or tokens that captures semantic meaning and relationships in a continuous vector space. Words with similar meanings have similar vector representations.

Feature Engineering: The process of manually selecting, extracting, and transforming raw data into features suitable for machine learning models.

Fine-tuning: The process of taking a pre-trained model and continuing its training on a specific downstream task with task-specific data, adapting the model's parameters to the new domain while retaining general language knowledge.

GloVe (Global Vectors for Word Representation): An unsupervised learning algorithm for obtaining word embeddings by aggregating global word-word co-occurrence statistics from a corpus, producing vectors where word similarity correlates with vector similarity.

Hyper parameter Tuning: The process of systematically searching for optimal hyper parameter values (e.g., learning rate, number of layers, regularization strength) that maximize model performance on validation data.

LIWC (Linguistic Inquiry and Word Count): A psycholinguistic dictionary and text analysis tool that categorizes words into psychological and linguistic dimensions, frequently used to extract features related to emotional and cognitive states.

LSTM (Long Short-Term Memory): A type of recurrent neural network designed to capture long-range dependencies in sequential data through

specialized gating mechanisms that control information flow.

MentalBERT: A domain-specific variant of BERT pre-trained on mental health-related text corpora, optimized for detecting mental health conditions from social media text.

N-gram: A contiguous sequence of n items (words, characters) from a text. For example, bigrams ($n=2$) from "mental health" would be "mental health". Used as features in text classification.

NLP (Natural Language Processing): A field of artificial intelligence focused on enabling computers to understand, interpret, and generate human language.

Overfitting: A modeling error occurring when a model learns training data too well, including noise and outliers, resulting in excellent training performance but poor generalization to new, unseen data.

Preprocessing: The initial data cleaning and transformation steps applied to raw text, including tokenization, normalization, stop word removal, and handling special characters.

RNN (Recurrent Neural Network): A class of neural networks designed for sequential data processing, where connections between nodes form directed cycles, allowing information to persist across time steps.

RoBERTa (Robustly Optimized BERT Pretraining Approach): An optimized version of BERT with modifications to the training procedure, including removing the next sentence prediction objective and training on larger batches with more data, resulting in improved performance.

Sentiment Analysis: The computational task of identifying and extracting subjective information, emotions, opinions, and attitudes from text, typically classifying text as positive, negative, or neutral.

Suicidal Ideation: Thoughts about, consideration of, or planning of suicide. Ranges from passive thoughts (e.g., "I wish I were dead") to active planning. Early detection is critical for intervention.

SVM (Support Vector Machine): A classical machine learning algorithm that finds an optimal hyperplane to separate classes in high-dimensional feature space, often using kernel functions (linear, polynomial, RBF) to handle non-linear relationships.

TF-IDF (Term Frequency-Inverse Document Frequency): A statistical measure that evaluates the importance of a word in a document relative to a collection of documents, commonly used for text feature representation.

Tokenization: The process of breaking text into smaller units (tokens) such as words, sub words, or characters for computational processing.

Transfer Learning: A machine learning technique where a model pre-trained on one task is fine-tuned and adapted for a related task, leveraging previously learned representations.

Transformer: A neural network architecture based entirely on attention mechanisms, processing entire sequences simultaneously rather than sequentially, enabling parallel computation and better capture of long-range dependencies.

Word2Vec: A technique for learning word embeddings by training a shallow neural network to predict context words (CBOW) or predict a word from its context (Skip-gram), producing dense vector representations that capture semantic relationships.

6. References

- [1] Y. Cao, J. Dai, Z. Wang, Y. Zhang, X. Shen, Y. Liu and Y. Tian, "Machine Learning Approaches for Mental Illness Detection on Social Media: A Systematic Review of Biases and Methodological Challenges," *arXiv*, 2025.
- [2] Z. Ding, Z. Wang, Y. Zhang, Y. Cao, Y. Liu, X. Shen, Y. Tian and J. Dai, "Trade-offs between machine learning and deep learning for mental illness detection on social media," *Scientific Reports*, vol. 15, 2025.
- [3] S. K. Sharma, A. I. Alutaibi, A. R. Khan, G. G. Tejani, F. Ahmad and S. J. Mousavirad, "Early detection of mental health disorders using machine learning models using behavioral and voice data analysis," *Scientific Reports*, vol. 15, 2025.
- [4] D. William and . D. Suhartono, "Text-based Depression Detection on Social Media Posts: A Systematic Literature Review," in *Proceedings of the 2020 International Conference on Computer Science and Computational Intelligence*, 2021.
- [5] B. G. Bokolo and Q. Liu, "Advanced Comparative Analysis of Machine Learning and Transformer Models for Depression and Suicide Detection in Social Media Texts," *Electronics*, vol. 13, 2024.
- [6] A. Senanayake and Z. Liang, "Anxiety Detection in Reddit Posts Through Emotion Dynamics," in *Proceedings of the 18th International Joint Conference on Biomedical Engineering Systems and Technologies (HEALTHINF 2025)*, 2025.
- [7] D. Owen, J. Camacho-Collados and L. Espinosa-Anke, "Towards Preemptive Detection of Depression and Anxiety in Twitter," in *Social Media Mining for Health Applications (SMM4H)*, 2020.
- [8] R. Safa, P. Bayat and L. Moghtader , "Automatic detection of depression symptoms in twitter using multimodal analysis," *Journal of Supercomputing*, vol. 78, pp. 4709-4744, 2022.
- [9] Z. Chen, R. Yang, S. Fu, N. Zong, H. Liu and M. Huang, "Detecting Reddit Users with Depression Using a Hybrid Neural Network SBERT-CNN," in *2023 IEEE 11th International Conference on Healthcare Informatics (ICHI 2023)*, 2023.
- [10] S. Zhou and M. Mohd, "Mental Health Safety and Depression Detection in Social Media Text Data: A Classification Approach Based on a Deep Learning Model," *IEEE Access*, vol. 13, pp. 63284 - 63297, 2025.
- [11] A. Anshul, G. S. Pranav, M. Z. Ur Rehman and N. Kumar, "A Multimodal Framework for Depression Detection During COVID-19 via Harvesting Social Media," *IEEE Transactions on Computational Social Systems*, vol. 11, pp. 2872 - 2888, 2023.
- [12] M. Mansoor and K. Ansari, "Early Detection of Mental Health Crises through Artificial-Intelligence-Powered Social Media Analysis: A Prospective Observational Study," *J Pers Med*, vol. 14, p. 958, 2024.
- [13] M. Kerasiotis, L. Ilias and D. Askounis, "Depression detection in social media posts using transformer-based models and auxiliary features," *Springer*, 2024.

- [14] B. G. Bokolo and Q. Liu, "Deep Learning-Based Depression Detection from Social Media: Comparative Evaluation of ML and Transformer Techniques," *Electronics*, vol. 12, p. 4396, 2023.
- [15] R. Chandrasekaran, S. Kotaki and A. H. Nagaraja, "Detecting and tracking depression through temporal topic modeling of tweets: insights from a 180-day study," *npj Mental Health Research*, vol. 3, 2024.
- [16] L. Bendebane, Z. Laboudi, A. Saighi, H. Al-Tarawneh, A. Ouannas and G. Grassi, "A Multi-Class Deep Learning Approach for Early Detection of Depressive and Anxiety Disorders Using Twitter Data," *Algorithms*, vol. 16, 2023.
- [17] N. Hussain, A. Qasim, G. Mehak, M. Zain, G. Sidorov, A. Gelbukh and O. Kolesnikova, "Multi-Level Depression Severity Detection with Deep Transformers and Enhanced Machine Learning Techniques," *AI*, vol. 6, p. 157, 2025.
- [18] N. M. Rum, N. A. Azura, F. A. Saharudin and A. Akbar, "Uncovering Depression on Social Media using BERT Model," *Journal of Advanced Research Design*, vol. 129, pp. 46-59, 2025.

7. Bibliography

- [1] Y. Cao, J. Dai, Z. Wang, Y. Zhang, X. Shen, Y. Liu and Y. Tian, "Machine Learning Approaches for Mental Illness Detection on Social Media: A Systematic Review of Biases and Methodological Challenges," *arXiv*, 2025.
- [2] Z. Ding, Z. Wang, Y. Zhang, Y. Cao, Y. Liu, X. Shen, Y. Tian and J. Dai, "Trade-offs between machine learning and deep learning for mental illness detection on social media," *Scientific Reports*, vol. 15, 2025.
- [3] S. K. Sharma, A. I. Alutaibi, A. R. Khan, G. G. Tejani, F. Ahmad and S. J. Mousavirad, "Early detection of mental health disorders using machine learning models using behavioral and voice data analysis," *Scientific Reports*, vol. 15, 2025.
- [4] D. William and . D. Suhartono, "Text-based Depression Detection on Social Media Posts: A Systematic Literature Review," in *Proceedings of the 2020 International Conference on Computer Science and Computational Intelligence*, 2021.
- [5] B. G. Bokolo and Q. Liu, "Advanced Comparative Analysis of Machine Learning and Transformer Models for Depression and Suicide Detection in Social Media Texts," *Electronics*, vol. 13, 2024.
- [6] A. Senanayake and Z. Liang, "Anxiety Detection in Reddit Posts Through Emotion Dynamics," in *Proceedings of the 18th International Joint Conference on Biomedical Engineering Systems and Technologies (HEALTHINF 2025)*, 2025.
- [7] D. Owen, J. Camacho-Collados and L. Espinosa-Anke, "Towards Preemptive Detection of Depression and Anxiety in Twitter," in *Social Media Mining for Health Applications (SMM4H)*, 2020.
- [8] R. Safa, P. Bayat and L. Moghtader , "Automatic detection of depression symptoms in twitter using multimodal analysis," *Journal of Supercomputing*, vol. 78, pp. 4709-4744, 2022.
- [9] Z. Chen, R. Yang, S. Fu, N. Zong, H. Liu and M. Huang, "Detecting Reddit Users with Depression Using a Hybrid Neural Network SBERT-CNN," in *2023 IEEE 11th International Conference on Healthcare Informatics (ICHI 2023)*, 2023.
- [10] S. Zhou and M. Mohd, "Mental Health Safety and Depression Detection in Social Media Text Data: A Classification Approach Based on a Deep Learning Model," *IEEE Access*, vol. 13, pp. 63284 - 63297, 2025.
- [11] A. Anshul, G. S. Pranav, M. Z. Ur Rehman and N. Kumar, "A Multimodal Framework for Depression Detection During COVID-19 via Harvesting Social Media," *IEEE Transactions on Computational Social Systems*, vol. 11, pp. 2872 - 2888, 2023.
- [12] M. Mansoor and K. Ansari, "Early Detection of Mental Health Crises through Artificial-Intelligence-Powered Social Media Analysis: A Prospective Observational Study," *J Pers Med*, vol. 14, p. 958, 2024.
- [13] M. Kerasiotis, L. Ilias and D. Askounis, "Depression detection in social media posts using transformer-based models and auxiliary features," *Springer*, 2024.

- [14] B. G. Bokolo and Q. Liu, "Deep Learning-Based Depression Detection from Social Media: Comparative Evaluation of ML and Transformer Techniques," *Electronics*, vol. 12, p. 4396, 2023.
- [15] R. Chandrasekaran, S. Kotaki and A. H. Nagaraja, "Detecting and tracking depression through temporal topic modeling of tweets: insights from a 180-day study," *npj Mental Health Research*, vol. 3, 2024.
- [16] L. Bendebane, Z. Laboudi, A. Saighi, H. Al-Tarawneh, A. Ouannas and G. Grassi, "A Multi-Class Deep Learning Approach for Early Detection of Depressive and Anxiety Disorders Using Twitter Data," *Algorithms*, vol. 16, 2023.
- [17] N. Hussain, A. Qasim, G. Mehak, M. Zain, G. Sidorov, A. Gelbukh and O. Kolesnikova, "Multi-Level Depression Severity Detection with Deep Transformers and Enhanced Machine Learning Techniques," *AI*, vol. 6, p. 157, 2025.
- [18] N. M. Rum, N. A. Azura, F. A. Saharudin and A. Akbar, "Uncovering Depression on Social Media using BERT Model," *Journal of Advanced Research Design*, vol. 129, pp. 46-59, 2025.
- [19] I. Tavchioski, M. Robnik-Šikonja and S. Pollak, "Detection of depression on social networks using transformers and ensembles," 2023.
- [20] M. M. Tadesse, H. Lin, B. Xu and L. Yang, "Detection of Suicide Ideation in Social Media Forums Using Deep Learning," *Algorithms*, vol. 1, 2019.
- [21] M. R. Islam, M. A. Kabir, A. Ahmed, A. R. M. Kamal, H. Wang and A. Ulhaq, "Depression detection from social network data using machine learning techniques," *Health Information Science and Systems*, vol. 6, 2018.
- [22] M. M. Tadesse, H. Lin, B. Xu and L. Yang, "Detection of Depression-Related Posts in Reddit Social Media Forum," *IEEE Access*, vol. 7, pp. 44883 - 44893, 2019.